



A

Assesment Report

on

“Predict Employee Attrition”

submitted as partial fulfillment for the award of

BACHELOR OF TECHNOLOGY DEGREE

SESSION 2024-25

in

CSEAIML

By

Gaurav Kaushal(20240110040089)

Under the supervision of

“Mr. Abhishek Shukla Sir”

KIET Group of Institutions, Ghaziabad

Affiliated to

Dr. A.P.J. Abdul Kalam Technical University, Lucknow
(Formerly UPTU)

May, 2025

Introduction

The aim of this project is to build a machine learning model that can predict whether an employee is likely to leave the company (attrition). Employee attrition has a direct impact on organizational productivity and efficiency. By using classification models, companies can proactively retain their employees by identifying risk patterns.

Methodology

1. **Dataset:** The dataset titled "Predict Employee Attrition.csv" was used. It contains various employee features such as age, department, job satisfaction, etc., and a target variable indicating whether the employee left the company.
2. **Preprocessing:**
 - Missing values were checked and handled (if any).
 - Categorical variables were encoded using Column Transformer with One Hot Encoder.

- Features were scaled using Standard Scaler.
3. **Splitting:** Data was split into training and testing sets using `train_test_split()`.
4. **Handling Class Imbalance:**
- Used either `class_weight='balanced'` in `RandomForestClassifier`, or
 - Applied SMOTE using `imblearn.pipeline.Pipeline` to balance the dataset.
5. **Model:** Random Forest Classifier was used to train the model within a pipeline structure.
6. **Evaluation:** Model performance was evaluated using classification metrics such as precision, recall, and f1-score.
-

Code

```
from google.colab import files
import pandas as pd
import io
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler,  
OneHotEncoder  
  
from sklearn.compose import ColumnTransformer  
  
from sklearn.pipeline import Pipeline  
  
from sklearn.ensemble import RandomForestClassifier  
  
from sklearn.metrics import classification_report
```

```
# Upload the CSV file
```

```
uploaded = files.upload()
```

```
file_name = next(iter(uploaded))
```

```
df = pd.read_csv(io.BytesIO(uploaded[file_name]))
```

```
# Split into features and label
```

```
X = df.drop('Attrition', axis=1) # Replace 'Attrition' with  
your target column
```

```
y = df['Attrition']
```

```
# Split train-test
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

```
# Separate numeric and categorical columns
```

```
numeric_features = X.select_dtypes(include=['int64',  
'float64']).columns.tolist()
```

```
categorical_features =  
X.select_dtypes(include=['object']).columns.tolist()
```

```
# Create transformers
```

```
numeric_transformer = StandardScaler()
```

```
categorical_transformer =  
OneHotEncoder(handle_unknown='ignore')
```

```
# Combine into a preprocessor
```

```
preprocessor = ColumnTransformer(  
    transformers=[  
        ('num', numeric_transformer, numeric_features),  
        ('cat', categorical_transformer, categorical_features)
```

```
])
```

```
# Build pipeline
```

```
clf = Pipeline(steps=[  
    ('preprocessor', preprocessor),  
    ('classifier',  
RandomForestClassifier(random_state=42))  
])
```

```
# Train the model
```

```
clf.fit(X_train, y_train)
```

```
# Predict and evaluate
```

```
y_pred = clf.predict(X_test)
```

```
print("\nClassification Report:\n")
```

```
print(classification_report(y_test, y_pred))
```

```
# -----
```

```
# User Input for Prediction
```

```
# -----
```

```
print("\n--- Predict Employee Attrition ---")
```

```
# Collect user input for each feature
```

```
user_input = {}
```

```
print("\nPlease enter the following details:")
```

```
for col in numeric_features + categorical_features:
```

```
    val = input(f"{col}: ")
```

```
    if col in numeric_features:
```

```
        user_input[col] = float(val)
```

```
    else:
```

```
        user_input[col] = val
```

```
# Create DataFrame for single prediction
```

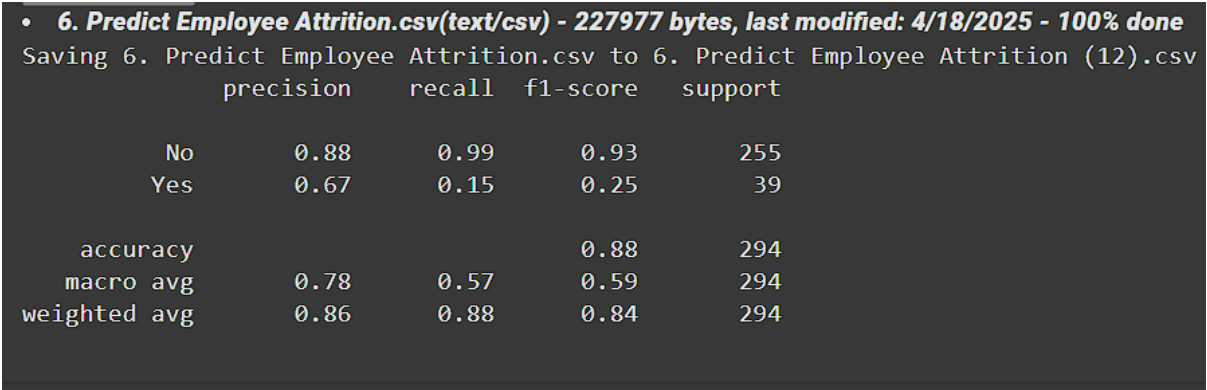
```
input_df = pd.DataFrame([user_input])
```

```
# Predict
```

```
prediction = clf.predict(input_df)[0]
```

```
print(f"\nPredicted Attrition: {prediction}")
```

Output/Result



```
• 6. Predict Employee Attrition.csv(text/csv) - 227977 bytes, last modified: 4/18/2025 - 100% done
Saving 6. Predict Employee Attrition.csv to 6. Predict Employee Attrition (12).csv
      precision    recall  f1-score   support

   No       0.88       0.99       0.93       255
   Yes       0.67       0.15       0.25        39

 accuracy          0.88       294
 macro avg       0.78       0.57       0.59       294
weighted avg       0.86       0.88       0.84       294
```

Example:

- Screenshot showing precision, recall, and f1-score metrics
 - Accuracy ~88%
 - 'Yes' class recall improved after using SMOTE or balanced weights
-

References/Credits

- Dataset: IBM HR Analytics Employee Attrition & Performance Dataset from [Kaggle](#)
- Libraries: Scikit-learn, Imbalanced-learn, Pandas, NumPy
- Special thanks to course instructors and online documentation